

SLR210

Project: Obstruction-Free Consensus and Paxos

The goal of this project is to get an initial experience in designing a fault-tolerant distributed system. Here we focus on a state-machine replicated system build atop a consensus abstraction.

1 Specification

An obstruction-free consensus (OFC) algorithm exports one operation *propose*(*v*) with an input value in a set $v \in V = \{0, 1\}$. When a process invokes *propose*(*v*), we say that the process *proposes* *v*. The operation returns either a value $v' \in V$ (in which case we say that the process *decides* v') or a special value *abort* $\notin V$ (in which case we say that the invocation *aborts*). A process can invoke the *propose* operation multiple times.

The following properties must be met:

- Validity: every decided value is a proposed value.
- Agreement: no two processes decide differently.
- Obstruction-free termination:
 - If a correct process proposes, it eventually decides or aborts.
 - If a correct process decides, no correct process aborts infinitely often.
 - If there is a time after which *exactly one* correct process *p* proposes a value sufficiently many times, *p* eventually decides.

2 Concurrent environment

The goal of the project is to implement OFC for the following environment:

- We have N asynchronous processes. Every process has a distinct *identifier*. The identifiers are publicly known.
- Every two processes can communicate via a reliable asynchronous point-to-point channel.
- Up to $f < N/2$ of the processes are subject to crash failures: a faulty process prematurely stops taking steps of its algorithm. A process that never crashes is called correct.

3 Prerequisites

The project assumes some rudimentary knowledge of Java and Maven. For starters, familiarize yourself with the Java version of AKKA, an actor-based programming model <https://akka.io/docs/>. In particular, we advise you to read the motivation behind the actor model and the general overview provided in the AKKA documentation. You may also read the wikipedia entry about actors.

Once you have read the above material, you may download the skeleton of the project. This skeleton is available as a tarball in the shared repository on Google Drive. If you are familiar with git, you may instead clone (or fork) the following repository.

4 Formalities

The project is pursued in teams of **two or three students**.

The implemented system should be provided with a short report. The team should also prepare a short presentation to be given at the end of the course.

The meeting on **04/03** serves as an introduction to the project. The reports should be submitted by **15/04**. To submit your report, send an email entitled “SLR 210 - your names” to Prof. Pierre Sutra, at the following address: *pierre.sutra@telecom-sudparis.eu*.

Project presentations are taking place on **04/04**.

5 Implementation

The implementation should extend the basic construction creating a system of a given size and ensure all-to-all connectivity. Create N actors (processes), and pass references of all N processes to each of them.

Use the name **Process** for the process class. For the *Process* class, create methods for invoking the operation *propose*, processing received messages, and returning response indications.

Our goal is to simulate runs of the algorithm in a system of N processes out of which up to f can fail. You can proceed as follows.

For every process, the **main** method sends a special *launch* message. Once process i receives a launch message, it picks an input value, randomly chosen in $\{0,1\}$ and invokes instances of *propose* operation with this value until a value is decided. As a basis, one can use the OFC pseudocode discussed in the lecture (adjusted to be used within AKKA).

The **main** method also selects f processes at random (e.g., using the shuffle method from `java.collections`) and sends each of them a special *crash* message. If a process receives a crash message it enters the *fault-prone* mode: for any processed *event* in the algorithm, the process decides, with a fixed probability α , if it going to *crash*. If it crashes, it enters the *silent* mode, not reacting to any future event.

Use the **LoggingAdapter** class to log both the timing of the invocation and the response of every operation each process performs.

- Emulate a leader election mechanism: after a fixed timeout t_{le} , the **main** method randomly picks up a process that is not *fault-prone* and sends a *hold* message to every other process. After receiving a *hold* message, a process stops invoking *propose* operations.

For example, by invoking `Thread.sleep(50)`, the **main** method “freezes” for 5ms.

An alternative method consists in using the scheduler. For example, the following command:

```
system.scheduler().scheduleOnce(Duration.create(50, TimeUnit.MILLISECONDS),
testActor, "foo", system.dispatcher(), null);
```

results in a message “foo” sent by the scheduler to **testActor** in 50ms.

- Perform the experiment for different system sizes N and leader election times t_{le} .

For each configuration, measure the *consensus latency*: how long it takes for the first process to decide.

Your goal is to find out how the latency depends on N (for a fixed t_{le}) and t_{le} (for a fixed N).

Also try to see if the probability of failures α (for fixed N and t_{le}) affects the latency.

- As a baseline, you can consider $N = 3, 10, 100$ (with $f = 1, 4, 49$, respectively), $\alpha = 0, 0.1, 1$.

Also you should select several values of t_{le} in a small enough range to capture the effect of leader election. Notice that if t_{le} is chosen to be too large, it is very likely that some process will decide before that due to a transient absence of contention.

So the right strategy would be to start with a large t_{le} and then gradually decrease it until no decision is observed.

Each experiments should be repeated 5 times and the average latency should be evaluated.

6 Report

Prepare a short report (up to 15 pages), preferably in English (can also be written in French if English does not feel comfortable). The report should contain:

- The statement of the problem that your implementation solves;
- A high level description of the implementation;
- A report on performance analysis, with plots relating the latency with N , t_{le} and α .

The report and the code of the implementation should be uploaded through the eCampus system by the designated deadline.

7 Presentation

The presentation (7 mins) should contain a brief overview of the main features of the algorithm, its correctness arguments and performance. We envision 10 minutes per team (including 3 minutes for questions), so the time bounds are strict.

The presentation and the report should be comprehensible for a non-expert in consensus protocols. Try to be pedagogical, imagine your peers as the audience. The grades depend on the clarity and completeness.